DATA-DRIVEN INSIGHTS INTO FOOD PRICE INSTABILITY

A Project Report
Submitted in the partial fulfillment of the
requirements for the award of the degree of

BACHELOR OF TECHNOLOGY

In

DEPARTMENT OF COMPUTER SCIENCE ENGINEERING

By

KAVYA KANNE - 2320030048

YANDAPALLI RUSHITA – 2320030091

NISHANTH BABBULA – 2320030119

Under the Esteemed Guidance of

LECTURE NAME
VINISHA VELTHURU

Koneru Lakshmaiah Education Foundation
(Deemed to be University estd. u/s. 3 of the UGC Act, 1956)
Off-Campus: Bachupally-Gandimaisamma Road, Bowrampet, Hyderabad, Telangana - 500 043.
Phone No: 7815926816, www.klh.edu.in

K L (Deemed to be) University

DEPARTMENT OF COMPUTER SCIENCE ENGINEERING

Declaration

The Project Report entitled **"Data driven insights into food price instability"** is a record of Bonafide work of team members **Kanne kavya–2320030048, Yandapalli Rushita – 2320030091, and Nishanth Babbula – 2320030119**, submitted in partial fulfilment for the award of B. Tech in Computer Engineering to K L University. The results embodied in this report have not been copied from any other departments/University/Institute.

Kanne Kavya - 2320030048

Yandapalli Rushita - 2320030091

Babbula Nishanth - 2320030119

K L (Deemed to be) University

DEPARTMENT OF COMPUTER SCIENCE ENGINEERING

Certificate

This is to certify that the project-based report entitled **"Data driven insights into food price instability"** is a Bonafide work carried out and submitted by **Kanne kavya–2320030048, Yandapalli Rushita – 2320030091, and Nishanth Babbula – 2320030119**, in partial fulfilment of the requirements for the award of the degree of **Bachelor of Technology in Computer Science Engineering**, at **K L (Deemed to be University)**, during the academic year **2025–2026**.

Signature of the Supervisor

Signature of the HOD **Signature of the External Examiner**

ACKNOWLEDGEMENT

The success in this project would not have been possible but for the timely help and guidance rendered by many people. Our wish to express my sincere thanks to all those who has assisted us in one way or the other for the completion of my project.

Our greatest appreciation to my guide, **Vinisha Velthuru**, Department of Computer Science which cannot be expressed in words for his tremendous support, encouragement and guidance for this project.

We express our gratitude to Dr.N. Sirisha, Head of the Department for Computer Science Engineering for providing us with adequate facilities, ways and means by which we can complete this project-based Lab.

We thank all the members of teaching and non-teaching staff members, and who have assisted me directly or indirectly for successful completion of this project. Finally, I sincerely thank my parents, friends and classmates for their kind help and cooperation during my work.

Kanne Kavya - 2320030048

Yandapalli Rushita - 2320030091

Babbula Nishanth - 2320030119

TABLE OF CONTENTS

S.No	Contents	Page no
1	Abstract	6
2	Introduction	7
3	Literature survey	9
4	Hardware and Software requirements	12
5	Implementation	16
6	Experimentation and Code	18
7	Results	21
8	Conclusion	24
19	References	25

ABSTRACT

Food price instability is a critical issue affecting millions of households in India, influencing food security, inflation, and farmer livelihoods. While financial markets are extensively analyzed using advanced tools, agricultural commodity prices often remain underexplored despite their significant socio-economic impact.

This project presents a data-driven Big Data analytics system to analyze, monitor, and predict food price volatility across Indian markets. The system leverages large-scale daily agricultural price data from the Agmarknet (Agricultural Marketing Information Network) platform and utilizes distributed technologies such as Apache Spark and Apache Hive for efficient data processing and analysis.

The proposed framework performs comprehensive time-series analysis, volatility measurement using statistical indicators, and anomaly detection using the Z-score method to identify sudden price spikes and irregular fluctuations. It also enables city-wise comparative analysis to understand regional market behavior and identify areas with high price instability.

To enhance usability, the system integrates an interactive dashboard that allows users to visualize trends, detect anomalies, and explore insights dynamically. Additionally, the framework supports integration with global benchmarks such as the FAO Food Price Index to provide a broader economic perspective.

The system is designed to be scalable, efficient, and adaptable for real-time monitoring. By transforming raw agricultural data into actionable insights, this project aims to assist policymakers, farmers, and market stakeholders in making informed decisions, reducing uncertainty, and improving the resilience of the food supply system.

INTRODUCTION

Food price instability is a significant challenge in developing economies like India, where a large portion of the population depends on agriculture for livelihood and food consumption. Fluctuations in the prices of essential agricultural commodities directly impact food security, inflation rates, and the economic stability of households. Sudden price spikes or drops can create uncertainty for farmers, traders, and consumers, making it difficult to plan production, distribution, and consumption effectively.

Traditionally, agricultural price analysis has been limited to basic statistical methods and periodic reports, which are often insufficient to capture real-time market dynamics. With the rapid growth of data generation in agricultural markets, particularly through platforms such as Agmarknet (Agricultural Marketing Information Network), there is an increasing need for advanced data-driven solutions that can process large-scale datasets efficiently and provide meaningful insights.

In this context, Big Data technologies play a crucial role in handling high-volume, high-velocity, and high-variety datasets. Tools such as Apache Spark and Apache Hive enable distributed data processing and scalable analytics, making them suitable for analyzing large agricultural datasets. These technologies allow for efficient computation of trends, patterns, and anomalies in commodity prices across different regions.

This project presents a comprehensive Food Price Volatility Analysis System that leverages Big Data analytics to study agricultural price behavior across Indian markets. The system processes large-scale daily commodity price data obtained from the Agmarknet platform and applies advanced analytical techniques such as time-series analysis, volatility measurement, and anomaly detection using statistical methods like Z-score. These techniques help in identifying sudden price spikes, irregular fluctuations, and emerging instability patterns.

Furthermore, the system provides city-wise comparative analysis to understand regional variations in price behavior, enabling the identification of markets that are more prone to volatility. An interactive dashboard is integrated into the system to visualize trends, anomalies, and key insights, allowing users to explore the data dynamically and make informed decisions.
Overall, this project aims to transform raw agricultural price data into actionable insights using modern data engineering and analytics techniques.

Problem Description and Objectives:

Food price instability is a critical issue in India, where fluctuations in agricultural commodity prices significantly impact food security, inflation, and the livelihoods of farmers and consumers. Prices of essential commodities often vary due to seasonal changes, supply-demand imbalances, and external factors such as weather conditions and market inefficiencies.

Although large volumes of agricultural price data are available through platforms like Agmarknet, existing systems are mostly limited to basic reporting and lack advanced analytical capabilities to extract meaningful insights. As a result, stakeholders face challenges in understanding price trends, identifying volatility patterns, and detecting sudden price spikes, which are essential for effective planning and decision-making.

The objective of this project is to develop a scalable and efficient Big Data-based system to analyze agricultural price data and address these challenges. The system focuses on preprocessing and transforming large datasets, performing time-series analysis to study price trends, and measuring volatility using statistical methods such as standard deviation and coefficient of variation. It also incorporates anomaly detection using the Z-score technique to identify sudden price spikes and abnormal fluctuations.

In addition, the system enables city-wise and commodity-wise comparative analysis to understand regional market behavior. An interactive dashboard is developed to visualize trends, volatility, and anomalies, allowing users to explore insights dynamically. Overall, the project aims to convert raw agricultural data into actionable insights, supporting better decision-making and contributing to a more stable and resilient food supply system.

LITERATURE SURVEY

The increasing availability of agricultural market data and the growing need for food security have made the analysis of commodity price fluctuations an important research area. In countries like India, where agriculture plays a crucial role in the economy, price instability can significantly affect farmers, consumers, and policymakers. Traditional data processing methods are often inadequate for handling the large volumes of data generated from agricultural markets. As a result, Big Data analytics has emerged as a key solution for analyzing agricultural price trends and market behavior efficiently.

Agricultural Price Analysis and Market Dynamics

Agricultural price analysis has traditionally been performed using statistical techniques such as moving averages, regression models, and time-series forecasting. These methods help in identifying trends, seasonal patterns, and long-term price behavior. However, they are often limited in their ability to handle large-scale and continuously growing datasets. Researchers have highlighted that agricultural markets are influenced by multiple dynamic factors such as weather conditions, supply-demand imbalances, transportation costs, and government policies. These complexities make it difficult for traditional analytical models to provide accurate and real-time insights.

Big Data Analytics in Agricultural Systems

With the advancement of Big Data technologies, researchers have increasingly focused on using distributed computing frameworks for agricultural data analysis. Technologies such as Apache Hadoop, Apache Spark, and Apache Hive enable the processing of large datasets in a scalable and efficient manner. Apache Spark, in particular, has gained popularity due to its in-memory processing capability, which significantly reduces computation time compared to traditional disk-based systems. Studies have shown that Spark-based analytics can handle high-volume agricultural datasets and perform complex transformations and aggregations efficiently.

Big Data platforms also support integration with multiple data sources, allowing the combination of structured and unstructured data for comprehensive analysis. This capability is essential in agricultural systems, where data is collected from various sources such as market records, weather data, and supply chain information. The use of distributed systems ensures high scalability, fault tolerance, and improved performance, making them suitable for real-time and large-scale agricultural analytics.

Volatility Analysis in Commodity Prices

Volatility analysis is a critical component in understanding price instability in agricultural markets. Researchers have widely used statistical measures such as standard deviation and coefficient of variation to quantify price fluctuations. These measures provide insights into the degree of variability in commodity prices over time and help identify commodities that exhibit high instability. Studies indicate that high volatility is often associated with increased market risk and uncertainty, which can negatively impact both producers and consumers.

In addition to basic statistical methods, some research has explored advanced time-series models for volatility analysis. However, these models often require complex parameter tuning and may not be suitable for real-time applications. As a result, simpler statistical measures combined with Big Data processing have been widely adopted for large-scale agricultural analysis.

Anomaly Detection in Time-Series Data

Anomaly detection has become an important area of research for identifying unusual patterns in data. In the context of agricultural prices, anomalies may represent sudden price spikes or drops caused by unexpected events such as supply disruptions or market shocks. Detecting these anomalies is essential for early warning systems and effective decision-making.

Researchers have applied various techniques for anomaly detection, including rule-based methods, statistical approaches, and machine learning models. Rule-based methods involve setting thresholds to identify abnormal values, while statistical techniques such as Z-score analysis detect outliers based on deviations from the mean. The Z-score method is widely used due to its simplicity, interpretability, and effectiveness in identifying extreme values in large datasets.

Recent studies have also explored machine learning approaches such as clustering and classification for anomaly detection. While these methods can improve accuracy, they often require large labeled datasets and higher computational resources. Therefore, statistical methods remain a practical and efficient choice for large-scale agricultural data analysis.

Data Visualization and Decision Support Systems

Data visualization plays a vital role in converting complex datasets into understandable and actionable insights. Researchers have emphasized the importance of visual analytics in supporting decision-making processes. Visualization techniques such as line charts, bar graphs, heatmaps, and scatter plots are commonly used to represent trends, comparisons, and anomalies in data.

Interactive dashboards have become an essential component of modern data analytics systems. Tools such as Plotly, Tableau, and Streamlit allow users to interact with data, apply filters, and explore insights dynamically. These dashboards enhance user experience and enable stakeholders to quickly interpret results and make informed decisions. Visualization is particularly important in agricultural analytics, where users may not have a technical background and require intuitive representations of data.

Comparative Analysis of Traditional and Big Data Approaches

Several studies comparing traditional data processing methods with Big Data-based systems have concluded that distributed computing frameworks offer significant advantages in terms of scalability, speed, and efficiency. Traditional systems often struggle with large datasets and require significant processing time, whereas Big Data technologies enable parallel processing and faster computation.

Research has also shown that integrating statistical analysis with Big Data platforms improves the accuracy and reliability of results. The combination of time-series analysis, volatility measurement, and anomaly detection provides a comprehensive approach to understanding price behavior. Furthermore, the addition of visualization tools enhances the interpretability of analytical results, making the system more effective for real-world applications.

Conclusion

The literature clearly indicates that Big Data analytics has become essential for analyzing large-scale agricultural datasets and understanding price volatility. From traditional statistical methods to modern distributed computing frameworks, research has consistently demonstrated the need for scalable and efficient analytical systems. Techniques such as volatility measurement and anomaly detection play a crucial role in identifying patterns and irregularities in price data, while visualization tools support effective decision-making.

This project builds upon these research findings by developing a comprehensive system that integrates Big Data technologies, statistical analysis, and interactive visualization for food price volatility analysis. By processing large-scale Agmarknet data, detecting anomalies using the Z-score method, and providing dynamic insights through a dashboard, the project addresses the limitations of existing approaches and demonstrates the practical application of modern data analytics techniques in the agricultural domain.

HARDWARE AND SOFTWARE REQUIREMENTS:

The Food Price Volatility Analysis System is designed to process and analyze large-scale agricultural data efficiently using a combination of local computing resources and scalable data processing frameworks. While the system can be extended to distributed environments using Big Data technologies such as Apache Spark and Hive, it can also be executed on a standard personal computer for development and visualization purposes.

Hardware Specifications

The project does not require highly specialized hardware, as most of the data processing can be optimized using efficient libraries and frameworks. However, a basic system configuration is necessary to ensure smooth execution of data preprocessing, analysis, and dashboard visualization.

Processor

A system with an Intel Core i3 processor or equivalent is sufficient to run the application. For better performance and faster data processing, an Intel Core i5 or higher is recommended. The processor is mainly responsible for executing data transformations, running analytical functions, and rendering visualizations in the dashboard.

RAM

A minimum of 4 GB RAM is required to run the application, including the dashboard interface and data processing scripts. However, 8 GB or more is recommended to handle larger datasets efficiently, especially when performing operations such as anomaly detection, aggregation, and visualization simultaneously.

Storage

The project requires minimal local storage as the dataset is stored in CSV format. Approximately 500 MB to 1 GB of free disk space is sufficient for storing the dataset, project files, and required libraries. Additional storage may be needed if larger datasets are used for extended analysis.

Graphics

No dedicated graphics card is required for this project. Integrated graphics available in modern processors are sufficient, as the system focuses on data analysis and visualization rather than intensive graphical computations.

OperatingSystem

The system is compatible with multiple operating systems, including Windows, Linux, and macOS. Since the project is developed using platform-independent tools such as Python and Streamlit, it can run seamlessly across different environments without any dependency issues.

Software Requirements

The project utilizes a combination of programming languages, frameworks, and tools to perform data processing, analysis, and visualization effectively.

ProgrammingLanguage
•Python (Version 3.x or above)
Python is used as the primary programming language due to its extensive support for data analysis, statistical computation, and visualization.

Libraries and Frameworks
The following Python libraries are used in the project:
• Pandas – For data cleaning, preprocessing, and manipulation
• NumPy – For numerical computations
• Plotly – For interactive data visualization
• Streamlit – For building the interactive dashboard
• FastAPI – For developing RESTful APIs
• SciPy – For statistical analysis and Z-score anomaly detection
These libraries provide efficient tools for handling large datasets, performing analytics, and presenting results in a user-friendly manner.

Big Data Technologies
• Apache Spark – For scalable data processing (optional/extended use)
• Apache Hive – For structured data querying (optional)
These technologies enhance the system's capability to process large datasets and support distributed computing environments.

Development Environment
• Visual Studio Code – For coding and debugging
• Jupyter Notebook (optional) – For experimentation and testing
• Web Browser (Google Chrome / Microsoft Edge) – For running the dashboard interface

Visualization Tools
• Streamlit Dashboard
• Plotly Charts (Line charts, Bar graphs, Heatmaps, Scatter plots)
These tools are used to visualize price trends, volatility patterns, anomalies, and comparative analysis across cities and commodities.

Summary
With the above hardware and software configuration, the Food Price Volatility Analysis System can be executed efficiently on a standard computing system. The project leverages powerful data analysis libraries and scalable technologies to process agricultural price data, detect volatility patterns, and provide interactive insights. The flexible design allows the system to run both on local machines and extended Big Data environments, ensuring scalability, performance, and ease of use.

IMPLEMENTATION

The implementation of the Food Price Volatility Analysis System is carried out using a modular and layered approach, integrating data preprocessing, analytics, backend services, and visualization components. The system is designed to efficiently process agricultural price data, perform statistical analysis, and present insights through an interactive dashboard.

Data Collection and Preprocessing

The system utilizes agricultural commodity price data obtained from the Agmarknet platform in CSV format. The dataset includes attributes such as commodity name, market (city), arrival date, and price details. During preprocessing, the raw data is cleaned by handling missing values, converting data types, and renaming columns for consistency. The date field is converted into a standard datetime format, and additional features such as year, month, and day are extracted to support time-series analysis. Furthermore, an average price column is computed using minimum, maximum, and modal prices to improve analytical accuracy.

Data Processing and Analytics

The core analytical operations are implemented in a dedicated analytics module. This module performs time-series aggregation, volatility measurement, and anomaly detection. Price trends are analyzed using daily, monthly, and yearly aggregations, along with moving averages to identify patterns over time. Volatility is measured using statistical indicators such as standard deviation and coefficient of variation, which help in understanding the degree of price fluctuation across commodities and cities.
Anomaly detection is implemented using the Z-score method, where each price value is compared against the mean and standard deviation of its respective commodity and city group. Values that exceed a predefined threshold are identified as anomalies, representing sudden price spikes or drops. This approach provides an efficient and interpretable way to detect irregular patterns in large datasets.

Backend API Development

A RESTful API is developed using FastAPI to provide structured access to the analytical results. The backend exposes multiple endpoints for retrieving price data, trends, volatility metrics, anomaly detection results, and city-wise comparisons. The dataset is loaded once and cached in memory to enhance performance and reduce redundant computations. The API also supports dynamic filtering based on commodity, city, and date range, allowing flexible data retrieval.

Interactive Dashboard Implementation

An interactive dashboard is developed using Streamlit to visualize the analytical outputs. The dashboard provides multiple sections, including price trends, volatility analysis, anomaly detection, seasonal patterns, and comparative analysis. Users can apply filters such as commodity selection, city selection, and date range to dynamically explore the data. Visualization is implemented using Plotly, which enables interactive charts such as line graphs for trend analysis, bar charts for comparisons, heatmaps for seasonal insights, and scatter plots for anomaly detection. Key performance indicators (KPIs) are displayed to summarize important metrics such as average price, number of anomalies, and most volatile commodity.

System Integration and Workflow

The system follows a structured workflow where the dataset is first preprocessed and stored in memory. The analytics module performs computations on the processed data, and the backend API provides access to these results. The Streamlit dashboard interacts with the analytics module or API to fetch processed data and display visualizations. This modular design ensures scalability, maintainability, and efficient execution.

Summary

The implementation integrates data preprocessing, statistical analysis, backend services, and interactive visualization into a unified system. By leveraging Python-based tools and scalable design principles, the system efficiently processes agricultural price data, detects volatility and anomalies, and presents meaningful insights through an intuitive interface.

EXPEREMENTATION AND CODE:

Code Structure

The experimentation phase of the Food Price Volatility Analysis System focuses on validating the effectiveness of data preprocessing, analytical techniques, and visualization methods used in the project. The system is tested using agricultural commodity price data obtained from the Agmarknet platform, which contains records across different cities, commodities, and time periods. Various experiments are conducted to analyze price trends, measure volatility, detect anomalies, and evaluate system performance.

Dataset and Experimental Setup

The dataset used for experimentation consists of large-scale agricultural price records stored in CSV format. The data includes attributes such as commodity name, market (city), arrival date, and pricing details including minimum, maximum, and modal prices. The dataset is loaded into the system using Python and processed using libraries such as Pandas and NumPy. The preprocessing stage ensures that the dataset is clean, consistent, and suitable for further analysis.

The experiments are conducted on a standard computing environment using Python-based tools. The system processes the dataset and generates outputs in the form of aggregated data, statistical metrics, and visual representations through the dashboard.

Data Preprocessing Implementation

The preprocessing step involves loading the dataset, renaming columns, handling missing values, and converting data types. The date field is transformed into a datetime format, and new features such as year, month, and average price are derived. These transformations improve data quality and enable efficient time-series analysis.

The implementation uses Pandas functions for data manipulation, ensuring that operations such as filtering, grouping, and aggregation are performed efficiently. The cleaned dataset is then used as input for further analytical computations.

Time-Series Analysis and Trend Evaluation

Experiments are conducted to analyze price trends over time using aggregation techniques. The system groups data based on different time intervals such as daily, monthly, and yearly periods. Moving averages are computed to smooth fluctuations and identify long-term patterns in price behavior.

The results show that time-series analysis helps in identifying seasonal trends and gradual changes in commodity prices. Visualization of these trends through line charts allows users to easily interpret market behavior over time.

Volatility Measurement

To evaluate price instability, experiments are performed using statistical measures such as standard deviation and coefficient of variation. These metrics are calculated for different commodities and cities to identify areas with high price fluctuations.

The results indicate that certain commodities exhibit higher volatility due to market dynamics, while others remain relatively stable. This analysis provides valuable insights into the risk associated with different agricultural products.

Anomaly Detection Implementation

Anomaly detection is implemented using the Z-score method to identify sudden price spikes and drops. For each commodity and city group, the mean and standard deviation are calculated, and the Z-score is computed for each data point. Values exceeding a predefined threshold are classified as anomalies.

Experimental results demonstrate that the Z-score method effectively detects unusual price behavior, including extreme spikes and drops. These anomalies are visualized using scatter plots, making it easier to identify irregular patterns in the dataset.

Backend and API Testing

The backend API developed using FastAPI is tested to ensure accurate and efficient data retrieval. Each endpoint is validated by sending requests with different parameters such as commodity, city, and date range. The responses are verified to ensure correctness and consistency.

Performance testing shows that caching the dataset improves response time and reduces computational overhead. The API successfully supports dynamic filtering and provides structured data for visualization.

Dashboard and Visualization Testing

The Streamlit dashboard is tested for functionality, usability, and performance. Different filters and input combinations are applied to verify that the dashboard updates dynamically and displays correct results. Various charts, including line graphs, bar charts, heatmaps, and scatter plots, are evaluated for clarity and accuracy.

The results confirm that the dashboard provides an intuitive interface for exploring data and understanding insights. Users can easily identify trends, compare markets, and detect anomalies through interactive visualizations.

RESULT

The results of the Food Price Volatility Analysis System demonstrate the effectiveness of the proposed approach in analyzing large-scale agricultural price data and extracting meaningful insights. The system successfully processes the dataset obtained from the Agmarknet platform and provides accurate analysis through statistical methods and interactive visualizations.

Trend Analysis Results

The time-series analysis performed on the dataset reveals clear patterns in commodity price behavior over different time periods. By aggregating data on daily, monthly, and yearly levels, the system identifies seasonal trends and long-term price movements. The use of moving averages helps in smoothing short-term fluctuations and highlighting underlying trends. The visual representation of these trends through line charts allows users to easily understand how prices vary over time.

Volatility Analysis Results

The system effectively measures price volatility using statistical metrics such as standard deviation and coefficient of variation. The results indicate that certain commodities and markets exhibit higher levels of fluctuation compared to others. This analysis helps in identifying unstable commodities and regions where price variations are more frequent. The comparative visualization of volatility across commodities and cities provides valuable insights into market stability.

Anomaly Detection Results

The implementation of anomaly detection using the Z-score method successfully identifies sudden price spikes and drops in the dataset. The system detects abnormal values by comparing each price with the mean and standard deviation of its respective group. The results show that the method is effective in highlighting extreme price variations, which are visualized using scatter plots and highlighted markers. These anomalies provide important information about unusual market behavior and potential disruptions.

Seasonal Analysis Results

The seasonal analysis reveals recurring patterns in commodity prices based on months and time periods. The system identifies specific months where prices tend to increase or decrease, indicating seasonal demand and supply effects. Heatmaps and bar charts are used to visualize these patterns, making it easier to interpret seasonal variations across different commodities.

City-wise Comparison Results

The system performs comparative analysis across different cities and markets to understand regional price behavior. The results show variations in average prices and volatility levels among different locations. Some cities exhibit higher price ranges and fluctuations, indicating regional differences in supply chains and market conditions. This comparison helps in identifying high-risk markets and understanding geographic trends.

Dashboard Output and User Interaction

The interactive dashboard provides a comprehensive view of all analytical results. Users can apply filters based on commodity, city, and date range to explore the dataset dynamically. The dashboard displays key performance indicators such as average price, total records, number of anomalies, and most volatile commodity. The combination of interactive charts and real-time updates enhances user experience and simplifies data interpretation.

Performance and Efficiency

The system demonstrates efficient performance in processing and analyzing the dataset. The use of optimized data processing techniques and caching mechanisms ensures faster execution and reduced computational overhead. The modular architecture allows smooth integration between data processing, backend services, and visualization components, resulting in a responsive and scalable system.

Summary

Overall, the results confirm that the Food Price Volatility Analysis System effectively analyzes agricultural price data, detects volatility patterns, and identifies anomalies. The integration of statistical methods and interactive visualization provides meaningful insights that can support decision-making for stakeholders. The system proves to be reliable, efficient, and suitable for real-world applications in agricultural market analysis.

CONCLUTION

The Food Price Volatility Analysis System successfully demonstrates the application of Big Data analytics and statistical techniques in understanding agricultural price behavior. The project effectively processes large-scale commodity price data obtained from the Agmarknet platform and transforms it into meaningful insights through systematic data preprocessing, analysis, and visualization.

The system integrates multiple analytical components, including time-series analysis, volatility measurement, and anomaly detection, to provide a comprehensive understanding of price fluctuations. The use of statistical methods such as standard deviation and coefficient of variation enables accurate measurement of price instability, while the implementation of the Z-score method effectively identifies sudden price spikes and abnormal variations. These techniques help in uncovering hidden patterns and irregularities in agricultural markets.

The development of an interactive dashboard enhances the usability of the system by allowing users to visualize trends, compare market behavior across cities, and explore data dynamically through filters and graphical representations. This makes the system not only technically efficient but also user-friendly and practical for real-world applications.

Overall, the project provides a scalable and efficient solution for analyzing food price volatility. It highlights the importance of data-driven decision-making in the agricultural sector and demonstrates how modern data analytics tools can be used to address real-world challenges. The system can serve as a valuable tool for policymakers, farmers, and market stakeholders in understanding market dynamics, reducing uncertainty, and improving the stability of the food supply system.

REFRENCES

- **References**

- **Agmarknet Portal**
  Government of India. Agricultural Marketing Information Network (Agmarknet).
  Available at: https://agmarknet.gov.in
  This reference was used as the primary data source for agricultural commodity prices across different markets in India.

- **Apache Spark Documentation**
  Apache Software Foundation. Apache Spark Documentation.
  Available at: https://spark.apache.org/docs/latest/
  This source provided understanding of distributed data processing and large-scale analytics used in the project.

- **Apache Hive Documentation**
  Apache Software Foundation. Apache Hive Documentation.
  Available at: https://hive.apache.org/
  This reference was used to understand data warehousing concepts and structured querying for large datasets.

- **Python Official Documentation**
  Python Software Foundation. Python Documentation.
  Available at: https://docs.python.org/3/
  Referenced for Python programming concepts, data handling, and implementation of analytical functions.

- **Anomaly Detection using Statistical Methods**
  Bolton, R. J., & Hand, D. J. (2002). Statistical Fraud Detection: A Review.
  This research provided insights into anomaly detection techniques such as identifying abnormal patterns using statistical methods like Z-score.

- **Big Data Analytics Concepts**
  Chen, M., Mao, S., & Liu, Y. (2014). Big Data: A Survey.
  This paper helped in understanding the role of Big Data analytics in processing large datasets efficiently.

- **Data Visualization Techniques**
  Few, S. (2009). Now You See It: Simple Visualization Techniques for Quantitative Analysis.
  This reference helped in understanding the importance of visual representation of data for better interpretation and decision-making.

- **Plotly Documentation**
  Plotly Technologies Inc. Plotly Documentation.
  Available at: https://plotly.com/python/
  Used for creating interactive visualizations such as line charts, bar charts, and heatmaps.